

From AI Usage to Workflow Rewiring

Uber's Agentic Pods show what an internal Agentic AI FDE model looks like when the objective is not tooling adoption, but business redesign.

99%

of engineers use AI tools

70%+

of PRs attributed to agents

2,500+

internal agent skills

16

Pods across 16 functions in 2 months

The real maturity signal is workflow redesign, not AI adoption.

Uber's numbers are impressive, but the deeper lesson is strategic: once engineering proves agent leverage internally, the next frontier is embedding AI builders into business workflows and redesigning those workflows end to end.

99%

Adoption is already solved

Near-universal internal AI usage means the organization has already crossed the experimentation threshold.

70%+

Agents are already in production work

PR attribution shows that AI is materially shaping delivery, not serving as a passive assistant.

2,500+

Internal skills compound

Reusable skills create a proprietary execution layer that new teams can inherit rather than rebuild.

10 days


Speed changes politics

Short-cycle shipping creates proof before skepticism or committee drag can kill momentum.

BOARD TAKEAWAY This is not a tooling story. It is a transformation model for discovering and scaling operational AI value inside the enterprise.

Agentic Pods compress discovery, build, validation, and shipping into one 10-day loop.

The brilliance of the operating model is that it treats workflow understanding as the first product milestone. Engineers are not handed tickets; they first learn how the work really happens.



WHAT MAKES THIS SUPERIOR TO A TOP-DOWN ROLLOUT


- It captures tacit knowledge that diagrams, SOPs, and vendor demos miss.
- It produces working proof before broad change-management machinery starts.
- It ties AI value to a named workflow owner, not a generic innovation budget.

FDE INTERPRETATION

- FDE becomes part field operator, part systems designer, part AI implementer.
- The scarce role is the bilingual builder who speaks domain nuance and agent mechanics.
- Pod cadence is fast enough to create a portfolio, not just isolated pilots.

The value comes from collapsing cycle time across complex workflows.

The strongest proof points are step-change reductions in elapsed time for real, messy, multi-system work.



The deepest lesson is that the workflow is now the unit of automation.

Discovery is the bottleneck

The best opportunities were hiding beside the operator, not inside existing documentation.

AI changes org design

The breakthrough comes from pairing AI-native engineers with domain-native experts in one execution unit.

Workflow redesign beats task speed

The biggest wins remove handoffs, approvals, tooling sprawl, and vendor dependence altogether.

INSIGHT 1

Discovery is the bottleneck.

The reason many enterprise AI programs stall is not model quality. It is that they try to automate from the outside in. Uber's pods start from observation, which converts invisible friction into buildable opportunities.

- Process diagrams capture the nominal path. Shadowing reveals the exception path and real decision logic.
-
- The best opportunities are usually attached to repeated friction and judgment bottlenecks that documentation understates.
-
- Agentic AI FDE should therefore begin with workflow immersion, not prompt brainstorming.

The deepest lesson is that the workflow is now the unit of automation.

Discovery is the bottleneck

The best opportunities were hiding beside the operator, not inside existing documentation.

AI changes org design

The breakthrough comes from pairing AI-native engineers with domain-native experts in one execution unit.

Workflow redesign beats task speed

The biggest wins remove handoffs, approvals, tooling sprawl, and vendor dependence altogether.

INSIGHT 2

AI changes org design.

The key innovation is not just technical. It is the pairing mechanism: one AI-proficient engineer with one domain expert creates a temporary but powerful bilingual operating cell.

- This pod structure collapses translation loss between the people who know the workflow and the people who can automate it.
- It suggests that internal AI transformation should be organized around pods, not around generic platform intake queues.
- The pod itself becomes the minimum viable transformation unit for an Agentic AI FDE team.

The deepest lesson is that the workflow is now the unit of automation.

Discovery is the bottleneck

The best opportunities were hiding beside the operator, not inside existing documentation.

AI changes org design

The breakthrough comes from pairing AI-native engineers with domain-native experts in one execution unit.

Workflow redesign beats task speed

The biggest wins remove handoffs, approvals, tooling sprawl, and vendor dependence altogether.

INSIGHT 3

Workflow redesign beats task speed.

The largest gains rarely come from making one step faster. They come from rebuilding the entire flow so that approvals, rework, tool switching, and redundant controls disappear.

- This is why reported time savings likely understate total value; the deeper savings include lower vendor spend and simpler operating architecture.
- Boards should ask whether AI is merely accelerating steps or actually collapsing structural friction across the workflow.
- The most strategic agent skills cross teams, tools, and organizational boundaries because workflows do too.

Agentic AI FDE should be treated as an embedded workflow transformation capability.

If an organization wants repeatable results, it needs a field execution model rather than a central AI help desk.

1

Embed

Put AI-native engineers next to the people doing the work so tacit knowledge becomes observable.

2

Diagnose

Map the workflow, exception paths, handoffs, data dependencies, and economics before choosing tooling.

3

Route

Match smaller models to repetitive operations and reserve frontier reasoning for complex judgment.

4

Govern

Control token cost, context reuse, approval boundaries, and risk classification from the start.

5

Validate

Prove the workflow improves for several practitioners, not just the original champion.

6

Scale

Turn each pod outcome into reusable playbooks, prompts, connectors, and model-task maps.

DISTILLED FDE POINT The winning internal AI team is not a generic center of excellence. It is a portfolio of embedded strike teams with codified reuse.

Different leaders should see different implications from the same operating model.

Board Operating model investment, not a software experiment. TOP PRIORITY Fund a dedicated Agentic AI FDE capability with portfolio accountability.	COO Use pods to remove friction from operational decision loops. TOP PRIORITY Choose the top workflows where time-to-decision damages service or growth.	CIO / CTO Orchestration discipline, not just model access. TOP PRIORITY Create a model-task map aligned to support, research, coding, ops, and docs.	CFO Measure AI value at workflow P&L level, not tool license level. TOP PRIORITY Build baseline economics before each pod sprint starts.
---	--	--	--

Different leaders should see different implications from the same operating model.

CFO

Measure AI value at workflow P&L level, not tool license level.

TOP PRIORITY

Build baseline economics before each pod sprint starts.

CFO Lens

AI value should be measured at workflow P&L level, not tool license level.

Finance leaders should look beyond labor substitution and quantify how workflow redesign changes cycle time, vendor spend, and decision quality.

- The largest economic wins may come from reducing rework, delays, and third-party process tooling.
- Token-cost discipline and model routing are direct margin levers.
- Self-service automation can shift capacity creation from headcount growth to workflow redesign.

Priority Moves

- Build baseline economics before each pod sprint starts.
- Separate local productivity gains from enterprise workflow gains.
- Reinvest savings into reusable infrastructure and high-value workflow expansion.

Different leaders should see different implications from the same operating model.

Board

Operating model investment, not a software experiment.

TOP PRIORITY

Fund a dedicated Agentic AI FDE capability with portfolio accountability.

Board Lens

This is an operating model investment, not a software experiment.

Boards should evaluate Agentic AI FDE as a mechanism to create proprietary workflow advantage, lower operating friction, and reshape decision velocity across functions.

- The moat is the internal workflow knowledge and reusable skill layer, not just model access.
- Capital allocation should favor workflows with measurable cross-functional impact rather than diffuse AI pilots.
- Governance should focus on proof of generalization, unit economics, and risk controls.

Priority Moves

- Fund a dedicated Agentic AI FDE capability with portfolio accountability.
- Track workflow cycle-time reduction, handoff elimination, and cost-to-serve gains.
- Require model-routing, token-cost, and risk-classification discipline from day one.

Different leaders should see different implications from the same operating model.

COO

Use pods to remove friction from operational decision loops.

TOP PRIORITY

Choose the top workflows where time-to-decision damages service or growth.

COO Lens

Use pods to remove friction from operational decision loops.

For operations leaders, the appeal is faster throughput with fewer manual handoffs and less dependence on brittle process ownership.

- The best targets are recurring workflows with high exception handling and high coordination cost.
- Operators should co-own validation because usability and trust determine adoption.
- Scaling should follow proven workflows, not blanket automation mandates.

Priority Moves

- Choose the top workflows where time-to-decision damages service or growth.
- Appoint named workflow owners for each pod sprint.
- Measure how many approvals, touches, and tool switches are removed.

Different leaders should see different implications from the same operating model.

CIO / CTO

**Orchestration discipline,
not just model access.**

TOP PRIORITY

Create a model-task map aligned to support, research, coding, ops, and docs.

CIO / CTO Lens

The technical challenge is orchestration discipline, not just model access.

Technology leaders should read this as a prompt to build the internal routing, connector, evaluation, and reuse layer that makes pods scalable and economical.

- Smaller models should absorb extraction, tagging, summarization, and repetitive operations.
- Frontier models should be reserved for reasoning-heavy, coding-heavy, and high-risk flows.
- Reusable skills, connectors, and evaluation harnesses are the compounding assets.

Priority Moves

- Create a model-task map aligned to support, research, coding, ops, and internal docs.
- Instrument token economics and context reuse before scale increases spend.
- Standardize agent controls, logging, and reusable workflow components.

Use these as the crispest board-level points in the room.

These are intentionally sharp, short, and presentation-friendly while still anchored in the deeper operating insight from the source text.

AI maturity starts when workflow redesign begins.

Adoption metrics matter, but enterprise advantage appears only when work is rebuilt around agents.

Enterprise advantage appears only when work is rebuilt around agents.

The best AI opportunities are invisible from the outside.

You find them by watching the work, not by reading the process map.

You find them by watching the work, not by reading the process map.

Ten-day shipping loops beat months of AI theater.

Short cycles create proof, urgency, and organizational belief.

Short cycles create proof, urgency, and organizational belief.

Generalization is the real quality gate.

If several practitioners cannot use it, it is still a pilot.

If several practitioners cannot use it, it is still a pilot.

Use these as the crispest board-level points in the room.

These are intentionally sharp, short, and presentation-friendly while still anchored in the deeper operating insight from the source text.

The workflow, not the task, is now the automation unit.

Isolated task speedups rarely produce the biggest financial or strategic returns.

Isolated task speedups rarely produce the biggest returns.

Pair AI builders with domain operators.

That pairing is the fastest way to convert tacit friction into working product.

That pairing is the fastest way to convert friction into working product.

Model routing is a margin lever.

The winners will not just use more agents; they will route complexity, token cost, and risk better.

Winners route complexity, token cost, and risk better.

Internal Agentic AI FDE becomes a transformation engine.

Its job is to discover, redesign, operationalize, and scale workflow advantage across the business.

It discovers, redesigns, operationalizes, and scales workflow advantage.

The next move is to institutionalize an internal Agentic AI FDE motion.

Days 0-30

- Select 3-5 workflows with scale, repetition, and decision impact.
-
- Assign AI-native builders and named domain experts.
-
- Define baseline cycle time, handoffs, error rate, and cost-to-serve.

Days 31-60

- Run pod sprints that shadow, diagnose, build, validate, and ship.
-
- Create a model-task map and token economics dashboard.
-
- Capture reusable prompts, connectors, and controls from each win.

Days 61-90

- Expand only the workflows that generalize across multiple operators.
-
- Rationalize tooling and vendor spend the new workflow makes obsolete.
-
- Stand up a durable Agentic AI FDE team with portfolio ownership.

BOARD QUESTIONS TO ASK

- Which workflows have the highest hidden friction and decision latency?
- Where are we funding AI experimentation without workflow proof?
- What portfolio metrics show wins generalize and compound?

RISKS TO CONTROL

- Scaling pilots without model-routing discipline or token controls.
- Automating tasks without redesigning approvals and handoffs.
- Centralizing AI away from the operator, losing discovery fidelity.

END-STATE AMBITION

- Every major workflow re-evaluated in light of agent capability.
- A reusable internal capability that compounds with each pod.
- A faster, cheaper, more intelligent operating system for the company.